

```
In [1]: import os
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import tensorflow_addons as tfa

from kaggle_datasets import KaggleDatasets
import matplotlib.pyplot as plt
import numpy as np

from PIL import Image

AUTOTUNE = tf.data.experimental.AUTOTUNE
```

Problem Statement

The goal of this exercise is to train a generative adversarial network (GAN) on images of Monet paintings, and then use that network to generate novel images in the style of Money using regular photographs as input. The training data set consists of 300 256 x 256 RGB images of Monet paintings. The required output is 7,000+ images in the style of Monet, using the photos provided as input. I will be using a CycleGAN for this task.

GANs in general are composed of two primary parts: a generator and a discriminator. The generator takes input, possibly noise but also images as in this case, and attempts to create the desired output. The discriminator acts as a binary classifier and compares the generator's output to determine whether it belongs to some set of "real", labeled examples. GANs are so called because each component works against the other, the generator trying to produce better output to fool the discriminator, and the the discriminator trying to get better at distinguishing between "real" and generated output.

CycleGANs are a subset of GANs which introduce an element of "cycle consistency loss". This is an additional loss function that quantifies the loss from transforming an image twice, that is, from "real" input, to generated output, and back to "real" again.

EDA

```
In [2]: GCS_PATH = KaggleDatasets().get_gcs_path()
```

```
In [3]: MONET_FILENAMES = tf.io.gfile.glob(str(GCS_PATH + '/monet_tfrec/*.tfrec'))
print('Monet TFRecord Files:', len(MONET_FILENAMES))

PHOTO_FILENAMES = tf.io.gfile.glob(str(GCS_PATH + '/photo_tfrec/*.tfrec'))
print('Photo TFRecord Files:', len(PHOTO_FILENAMES))
```

Monet TFRecord Files: 5
Photo TFRecord Files: 20

In [4]: IMAGE_SIZE = [256, 256]

```
def decode_image(image):
    image = tf.image.decode_jpeg(image, channels=3)
    image = (tf.cast(image, tf.float32) / 127.5) - 1
    image = tf.reshape(image, [*IMAGE_SIZE, 3])
    return image

def read_tfrecord(example):
    tfrecord_format = {
        "image_name": tf.io.FixedLenFeature([], tf.string),
        "image": tf.io.FixedLenFeature([], tf.string),
        "target": tf.io.FixedLenFeature([], tf.string)
    }
    example = tf.io.parse_single_example(example, tfrecord_format)
    image = decode_image(example['image'])
    return image
```

In [5]:

```
def load_dataset(filenamees, labeled=True, ordered=False):
    dataset = tf.data.TFRecordDataset(filenamees)
    dataset = dataset.map(read_tfrecord, num_parallel_calls=AUTOTUNE)
    return dataset
```

In [6]:

```
monet_ds = load_dataset(MONET_FILENAMES, labeled=True).batch(1)
photo_ds = load_dataset(PHOTO_FILENAMES, labeled=True).batch(1)
```

In [7]:

```
def print_imgs(directory):
    # Get the list of image file names in the directory
    image_files = os.listdir(directory)[:9]

    # Create a figure and subplot grid
    fig, axes = plt.subplots(3, 3, figsize=(10, 10))

    # Iterate over the image files and plot them in the grid
    for i, image_file in enumerate(image_files):
        # Open the image using PIL
        image_path = os.path.join(directory, image_file)
        image = Image.open(image_path)

        # Plot the image on the corresponding subplot
        row = i // 3
        col = i % 3
        axes[row, col].imshow(image)
        axes[row, col].axis('off')

    # Adjust spacing between subplots
    plt.tight_layout()

    # Display the grid of images
    plt.show()
```

Example Monets

```
In [8]: print_imgs('/kaggle/input/monet-gan-getting-started/monet_jpg')
```



Example Photos

```
In [9]: print_imgs('/kaggle/input/monet-gan-getting-started/photo_jpg')
```




Helper Functions

In [10]: `OUTPUT_CHANNELS = 3`

```
def downsample(filters, size, apply_instancenorm=True):
    initializer = tf.random_normal_initializer(0., 0.02)
    gamma_init = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)

    result = keras.Sequential()
    result.add(layers.Conv2D(filters, size, strides=2, padding='same',
                             kernel_initializer=initializer, use_bias=False))

    if apply_instancenorm:
        result.add(tf.layers.InstanceNormalization(gamma_initializer=gamma_init))

    result.add(layers.LeakyReLU())
```

```
return result
```

```
In [11]: def upsample(filters, size, apply_dropout=False):
          initializer = tf.random_normal_initializer(0., 0.02)
          gamma_init = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)

          result = keras.Sequential()
          result.add(layers.Conv2DTranspose(filters, size, strides=2,
                                             padding='same',
                                             kernel_initializer=initializer,
                                             use_bias=False))

          result.add(tfa.layers.InstanceNormalization(gamma_initializer=gamma_init))

          if apply_dropout:
              result.add(layers.Dropout(0.5))

          result.add(layers.ReLU())

          return result
```

Modeling

I will be using a CycleGAN to create a model that can generate images in the style of Monet. It will use both down sampling and up sampling to transform the input image data.

Downsampling is useful as it acts as feature reduction, lowering the computational demands of the model. It also allows the model to capture broad details about the image. It can also help the model achieve translational invariance, which means that features can be recognized regardless of their position within an image.

Up sampling is useful because it allows the model to capture small features of the image. It can also help with properly aligning image features in space, and avoiding some kinds of artifacts, like checkerboarding.

```
In [12]: def Generator():
          inputs = layers.Input(shape=[256,256,3])

          # bs = batch size
          down_stack = [
              downsample(64, 4, apply_instancenorm=False), # (bs, 128, 128, 64)
              downsample(128, 4), # (bs, 64, 64, 128)
              downsample(256, 4), # (bs, 32, 32, 256)
              downsample(512, 4), # (bs, 16, 16, 512)
              downsample(512, 4), # (bs, 8, 8, 512)
              downsample(512, 4), # (bs, 4, 4, 512)
              downsample(512, 4), # (bs, 2, 2, 512)
              downsample(512, 4), # (bs, 1, 1, 512)
          ]
```

```

up_stack = [
    upsample(512, 4, apply_dropout=True), # (bs, 2, 2, 1024)
    upsample(512, 4, apply_dropout=True), # (bs, 4, 4, 1024)
    upsample(512, 4, apply_dropout=True), # (bs, 8, 8, 1024)
    upsample(512, 4), # (bs, 16, 16, 1024)
    upsample(256, 4), # (bs, 32, 32, 512)
    upsample(128, 4), # (bs, 64, 64, 256)
    upsample(64, 4), # (bs, 128, 128, 128)
]

initializer = tf.random_normal_initializer(0., 0.02)
last = layers.Conv2DTranspose(OUTPUT_CHANNELS, 4,
                              strides=2,
                              padding='same',
                              kernel_initializer=initializer,
                              activation='tanh') # (bs, 256, 256, 3)

x = inputs

# Downsampling through the model
skips = []
for down in down_stack:
    x = down(x)
    skips.append(x)

skips = reversed(skips[:-1])

# Upsampling and establishing the skip connections
for up, skip in zip(up_stack, skips):
    x = up(x)
    x = layers.Concatenate()([x, skip])

x = last(x)

return keras.Model(inputs=inputs, outputs=x)

```

```

In [13]: def Discriminator():
    initializer = tf.random_normal_initializer(0., 0.02)
    gamma_init = keras.initializers.RandomNormal(mean=0.0, stddev=0.02)

    inp = layers.Input(shape=[256, 256, 3], name='input_image')

    x = inp

    down1 = downsample(64, 4, False)(x) # (bs, 128, 128, 64)
    down2 = downsample(128, 4)(down1) # (bs, 64, 64, 128)
    down3 = downsample(256, 4)(down2) # (bs, 32, 32, 256)

    zero_pad1 = layers.ZeroPadding2D()(down3) # (bs, 34, 34, 256)
    conv = layers.Conv2D(512, 4, strides=1,
                        kernel_initializer=initializer,
                        use_bias=False)(zero_pad1) # (bs, 31, 31, 512)

    norm1 = tf.layers.InstanceNormalization(gamma_initializer=gamma_init)(conv)

    leaky_relu = layers.LeakyReLU()(norm1)

```

```

zero_pad2 = layers.ZeroPadding2D()(leaky_relu) # (bs, 33, 33, 512)

last = layers.Conv2D(1, 4, strides=1,
                    kernel_initializer=initializer)(zero_pad2) # (bs, 30, 30,

return tf.keras.Model(inputs=inp, outputs=last)

```

```

In [14]: monet_generator = Generator()
photo_generator = Generator()

monet_discriminator = Discriminator()
photo_discriminator = Discriminator()

```

```

In [15]: class CycleGan(keras.Model):
    def __init__(
        self,
        monet_generator,
        photo_generator,
        monet_discriminator,
        photo_discriminator,
        lambda_cycle=10,
    ):
        super(CycleGan, self).__init__()
        self.m_gen = monet_generator
        self.p_gen = photo_generator
        self.m_disc = monet_discriminator
        self.p_disc = photo_discriminator
        self.lambda_cycle = lambda_cycle

    def compile(
        self,
        m_gen_optimizer,
        p_gen_optimizer,
        m_disc_optimizer,
        p_disc_optimizer,
        gen_loss_fn,
        disc_loss_fn,
        cycle_loss_fn,
        identity_loss_fn
    ):
        super(CycleGan, self).compile()
        self.m_gen_optimizer = m_gen_optimizer
        self.p_gen_optimizer = p_gen_optimizer
        self.m_disc_optimizer = m_disc_optimizer
        self.p_disc_optimizer = p_disc_optimizer
        self.gen_loss_fn = gen_loss_fn
        self.disc_loss_fn = disc_loss_fn
        self.cycle_loss_fn = cycle_loss_fn
        self.identity_loss_fn = identity_loss_fn

    def train_step(self, batch_data):
        real_monet, real_photo = batch_data

        with tf.GradientTape(persistent=True) as tape:

```

[illegible]


```

self.p_disc_optimizer.apply_gradients(zip(photo_discriminator_gradients,
                                         self.p_disc.trainable_variables))

return {
    "monet_gen_loss": total_monet_gen_loss,
    "photo_gen_loss": total_photo_gen_loss,
    "monet_disc_loss": monet_disc_loss,
    "photo_disc_loss": photo_disc_loss
}

```

Define loss functions

The loss functions below are how the model determines which weights need to be updated, and with what strength and direction. The cycle loss function is the distinguishing feature of CycleGANs, and it encourages bi-directional stability of the transformation. That is, an image that is converted should be able to be converted back and appear close to the original.

```

In [16]: def discriminator_loss(real, generated):
         real_loss = tf.keras.losses.BinaryCrossentropy(from_logits=True, reduction=tf.k

         generated_loss = tf.keras.losses.BinaryCrossentropy(from_logits=True, reduction

         total_disc_loss = real_loss + generated_loss

         return total_disc_loss * 0.5

```

```

In [17]: def generator_loss(generated):
         return tf.keras.losses.BinaryCrossentropy(from_logits=True, reduction=tf.keras.

```

```

In [18]: def calc_cycle_loss(real_image, cycled_image, LAMBDA):
         loss1 = tf.reduce_mean(tf.abs(real_image - cycled_image))

         return LAMBDA * loss1

```

```

In [19]: def identity_loss(real_image, same_image, LAMBDA):
         loss = tf.reduce_mean(tf.abs(real_image - same_image))
         return LAMBDA * 0.5 * loss

```

Train Model

```

In [20]: monet_generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
         photo_generator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)

         monet_discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)
         photo_discriminator_optimizer = tf.keras.optimizers.Adam(2e-4, beta_1=0.5)

```

```

In [21]: cycle_gan_model = CycleGan(
         monet_generator, photo_generator, monet_discriminator, photo_discriminator
         )

```

```

cycle_gan_model.compile(
    m_gen_optimizer = monet_generator_optimizer,
    p_gen_optimizer = photo_generator_optimizer,
    m_disc_optimizer = monet_discriminator_optimizer,
    p_disc_optimizer = photo_discriminator_optimizer,
    gen_loss_fn = generator_loss,
    disc_loss_fn = discriminator_loss,
    cycle_loss_fn = calc_cycle_loss,
    identity_loss_fn = identity_loss
)

```

```

In [22]: cycle_gan_model.fit(
    tf.data.Dataset.zip((monet_ds, photo_ds)),
    epochs=10
)

```

```

Epoch 1/10
300/300 [=====] - 188s 443ms/step - monet_gen_loss: 5.1248
- photo_gen_loss: 5.3362 - monet_disc_loss: 0.6581 - photo_disc_loss: 0.6323
Epoch 2/10
300/300 [=====] - 136s 452ms/step - monet_gen_loss: 3.6282
- photo_gen_loss: 3.7505 - monet_disc_loss: 0.6608 - photo_disc_loss: 0.6269
Epoch 3/10
300/300 [=====] - 136s 450ms/step - monet_gen_loss: 3.5729
- photo_gen_loss: 3.7091 - monet_disc_loss: 0.6514 - photo_disc_loss: 0.6208
Epoch 4/10
300/300 [=====] - 135s 450ms/step - monet_gen_loss: 3.4259
- photo_gen_loss: 3.5794 - monet_disc_loss: 0.6454 - photo_disc_loss: 0.6067
Epoch 5/10
300/300 [=====] - 135s 450ms/step - monet_gen_loss: 3.2899
- photo_gen_loss: 3.4240 - monet_disc_loss: 0.6340 - photo_disc_loss: 0.6087
Epoch 6/10
300/300 [=====] - 135s 450ms/step - monet_gen_loss: 3.2282
- photo_gen_loss: 3.3212 - monet_disc_loss: 0.6202 - photo_disc_loss: 0.6074
Epoch 7/10
300/300 [=====] - 135s 450ms/step - monet_gen_loss: 3.2189
- photo_gen_loss: 3.2947 - monet_disc_loss: 0.6179 - photo_disc_loss: 0.6124
Epoch 8/10
300/300 [=====] - 135s 449ms/step - monet_gen_loss: 3.2154
- photo_gen_loss: 3.2573 - monet_disc_loss: 0.6098 - photo_disc_loss: 0.6141
Epoch 9/10
300/300 [=====] - 135s 449ms/step - monet_gen_loss: 3.1956
- photo_gen_loss: 3.2087 - monet_disc_loss: 0.6068 - photo_disc_loss: 0.6199
Epoch 10/10
300/300 [=====] - 135s 449ms/step - monet_gen_loss: 3.1709
- photo_gen_loss: 3.1810 - monet_disc_loss: 0.6047 - photo_disc_loss: 0.6138

```

```

Out[22]: <keras.callbacks.History at 0x78a554094290>

```

Example Generated Images

```

In [23]: _, ax = plt.subplots(5, 2, figsize=(12, 12))
    for i, img in enumerate(photo_ds.take(5)):

```

```
prediction = monet_generator(img, training=False)[0].numpy()
prediction = (prediction * 127.5 + 127.5).astype(np.uint8)
img = (img[0] * 127.5 + 127.5).numpy().astype(np.uint8)

ax[i, 0].imshow(img)
ax[i, 1].imshow(prediction)
ax[i, 0].set_title("Input Photo")
ax[i, 1].set_title("Monet-esque")
ax[i, 0].axis("off")
ax[i, 1].axis("off")

plt.tight_layout()
plt.show()
```

Input Photo



Monet-esque



Input Photo



Monet-esque



Input Photo



Monet-esque



Input Photo




Monet-esque



Input Photo



Monet-esque



Conclusion

The generated images are not radically different from the input images. There is some indication that they are trending towards a more impressionistic and less photorealistic appearance, but it is clear they are not quite there. Unfortunately, as this is an extremely resource-intensive model, my ability to experiment with different parameters is limited. I suspect that increasing training time would significantly improve the results, as would more elaborate up and downsampling functions.

Submission

```
In [24]: import PIL
! mkdir '/kaggle/working/images'
```

```
In [ ]: i = 1
for img in photo_ds:
    prediction = monet_generator(img, training=False)[0].numpy()
    prediction = (prediction * 127.5 + 127.5).astype(np.uint8)
    im = PIL.Image.fromarray(prediction)
    im.save("/kaggle/working/images/" + str(i) + ".jpg")
    i += 1
```

```
In [ ]: import shutil
shutil.make_archive("/kaggle/working/images", 'zip', "/kaggle/working")
```

References

https://keras.io/examples/generative/dcgan_overriding_train_step/

<https://www.kaggle.com/code/amyjang/monet-cyclegan-tutorial>

<https://www.kaggle.com/c/gan-getting-started/data>

Github

<https://github.com/obbrown1/Deep-Learning-Week-5>